

# Basics of Machine Learning

TSIA-SD 210 - P3

Lecture 6 - Ensemble methods: bagging, random forests, boosting

---

Florence d'Alché-Buc

Contact: [florence.dalche@telecom-paristech.fr](mailto:florence.dalche@telecom-paristech.fr),  
Télécom ParisTech, France

# Table of contents

1. Motivation
2. Bagging
3. Random forests
4. Boosting
5. References

Motivation

Bagging

Random forests

Boosting

References

# Ensemble methods for classification and regression

1. Remark:
  - Machine Learning not so "automatic": too many hyperparameters to tune
2. **meta-learning**: a procedure that learns to learn
3. **committee learning** or **wisdom of the crowd**: better results are obtained by combining the predictions of a set of **diverse** classifiers/regressors
4. **ensemble learning**: Improve upon a single base predictive model by building an ensemble of predictive model (with no hyperparameter)

# Ensemble methods for regression

Let  $f_t, t = 1, \dots, T$  be  $T$  different regressors.

Notations:

$$\begin{aligned}\epsilon_t(x) &= y - f_t(x) \\ MSE(f_t) &= \mathbb{E}[\epsilon_t(x)^2] \\ f_{ens}(x) &= \frac{1}{T} \sum_t f_t(x) \\ &= y - \frac{1}{T} \sum_t \epsilon_t(x).\end{aligned}$$

## Encourage the diversity of base models

$$MSE(f_{ens}) = \mathbb{E}[(y - f_{ens}(x))^2]$$

If  $\epsilon_t$  are mutually independent with zero mean, then we have:

$$MSE(f_{ens}) = \frac{1}{T^2} \mathbb{E}[\sum_t \epsilon_t(x)^2]$$

The more diverse are the models, the more we reduce the mean square error !

# Ensemble methods for supervised classification

## Binary classification

$$h_{ens}(x) = \text{sign}\left(\sum_t h_t(x)\right)$$

## Multiclass classification

$$h_{ens}(x) = \arg \max_c \text{vote}(c, h_1, \dots, h_T)$$

$$\text{with : } \text{vote}(c, h_1, \dots, h_T) = \sum_t 1_{h_t(x)=c}(h_t(x))$$

- **Encourage the diversity of base predictors by:**
  - using bootstrap samples (Bagging and Random forests)
  - using randomized predictors (ex: Random forests)
  - using weighted version of the current sample (Boosting) with weights dependent on the previous predictor (adaptive sampling)



# Ensemble methods at a glance

- 1995: Boosting, Freund and Schapire
- 1996: Bagging, Breiman
- 2001: Random forests, Breiman
- 2006: Extra-trees, Geurts, Ernst, Wehenkel

Motivation

**Bagging**

Random forests

Boosting

References

## Reminder: Decomposition bias/variance in regression

Given  $x$ ,

$$\mathbb{E}_S \mathbb{E}_{y|x} (y - f_S(x))^2 = \text{noise}(x) + \text{bias}^2(x) + \text{variance}(x) \quad (1)$$

$\text{noise}(x)$ :  $\mathbb{E}_{y|x} [(y - E_{y|x}(y))^2]$ :

quantifies the error made by the Bayes model ( $E_{y|x}(y)$ )

$\text{bias}^2(x) = (E_{y|x}(y) - E_S[f_S(x)])^2$

measures the difference between minimal error (Bayes error) and the average model

$\text{variance}(x) = E_S [(f_S(x) - E_S[f_S(x)])^2]$

measures how much  $h_S(x)$  varies from one training set to another

# Introduction to bagging (regression) - 1

Assume we can generate several training independent samples  $\mathcal{S}_1, \dots, \mathcal{S}_T$  from  $P(x, y)$ .

A first algorithm:

- draw  $T$  training independent samples  $\{\mathcal{S}_1, \dots, \mathcal{S}_T\}$
- learn a model  $f_t \in \mathcal{F}$  from each training sample  $\mathcal{S}_t; t = 1, \dots, T$
- compute the average model :  $f_{ens}(x) = \frac{1}{T} \sum_{t=1}^T f_t(x)$

## Introduction to bagging - 2

The bias ( $E_{S_1, \dots, S_T}[f_{ens}(x)] - f_{target}(x)$ ) remains the same because :

$$E_{S_1, \dots, S_T}[f_{ens}(x)] = \frac{1}{T} \sum_t E_{S_t}[f_t(x)] = E_S[f_S(x)]$$

But the variance is divided by T:

$$E_{S_1, \dots, S_T}[(f_{ens}(x) - E_{S_1, \dots, S_T}[f_{ens}(x)])^2] = \frac{1}{T} E_S[(f_S(x) - E_S[f_S(x)])^2]$$

**When is it useful?** When the learning algorithm is unstable, producing high variance estimators such as trees !

# Bagging (Breiman 1996)

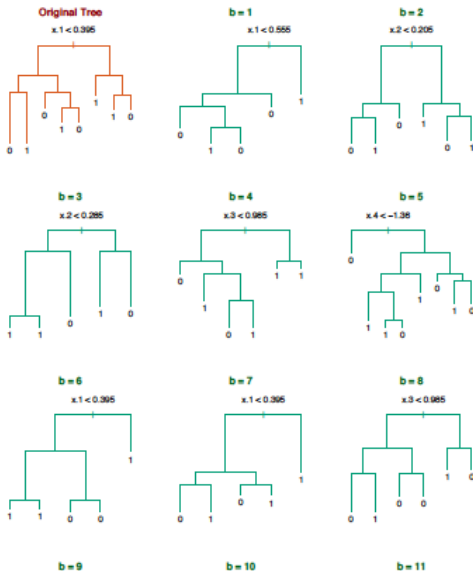
In practice, we do not know  $P(x,y)$  and we have only **one training sample**  $\mathcal{S}$ : we are going to use Bootstrap samples !

## Bagging = Bootstrap Aggregating

- draw  $T$  bootstrap samples  $\{\mathcal{S}_1, \dots, \mathcal{S}_T\}$  from  $\mathcal{S}$  (bootstrap: uniform sampling with replacement)
- Learn a model  $f_t$  for each  $\mathcal{S}_t$
- Build the average model:  $f_{bag}(x) = \frac{1}{T} \sum_t f_t(x)$


# Example of bagged trees

[Book: The elements of statistical learning, Hastie, Tibshirani, Friedman,



# Example of bagged trees

[Book: The elements of statistical learning, Hastie, Tibshirani, Friedman, 2001]





## Bagging in practise

- Variance is reduced but the bias can increase a bit (the effective size of a bootstrap sample is 30% smaller than the original training set  $\mathcal{S}$ )
- The obtained model is however more complex than a single model
- Bagging works for unstable predictors (neural nets, trees)
- In supervised classification, bagging a good classifier usually makes it better but bagging a bad classifier can make it worse

Motivation

Bagging

Random forests

Boosting

References

- Perturbe and combine algorithms
  - Perturbe the base predictive model
  - Combine the perturbed predictive model

REFS: Random forests: Breiman 2001

Geurts, Ernst, Wehenkel, Extra-trees, 2006

## Random forests algorithm

- INPUT:  $F = p$  candidate feature splits,  $\mathcal{S}_{train}$
- for  $t=1$  to  $T$ 
  - $\mathcal{S}_{train}^{(t)}$  m instance randomly drawn with replacement from  $\mathcal{S}_{train}$
  - $h_{tree}^{(t)} \leftarrow$  randomized decision tree learned from  $\mathcal{S}_{train}^{(t)}$
- OUTPUT:  $H^T = \frac{1}{T} \sum_t h_{tree}^{(t)}$

# Learning a single randomized tree

- To select a split at a node:
  - $R_f(F) \leftarrow$  randomly select (without replacement)  $f$  feature splits from  $F$  with  $f \ll p$
  - Choose the best split in  $R_f(F)$  (consider the different cut-points)
- Do not prune this tree

## Extra-trees

- INPUT: candidate feature splits  $F = \{1, \dots, p\}$ ,  $\mathcal{S}_{train}$
- for  $t=1$  to  $T$ 
  - Always use  $\mathcal{S}_{train}$
  - $h_{tree}^{(t)} \rightarrow$  : randomized decision tree learned from  $\mathcal{S}_{train}$
- OUTPUT:  $H^T = \frac{1}{T} h_{tree}^{(t)}$

# Extra-trees: learning a single randomized tree in extra-trees


To select a split at a node:

- randomly select (without replacement)  $K$  feature splits from  $F$  with  $K \ll |F|$
- Draw  $K$  splits using the procedure  $\text{Pick-a-random-split}(\mathcal{S}, i)$ :
  - let  $a_{\max}^i$  and  $a_{\min}^i$  denote the maximal and minimal value of  $x_i$  in  $\mathcal{S}$
  - Draw uniformly a cut-point  $a_c$  in  $[a_{\max}^i, a_{\min}^i]$
- Choose the best split among the  $K$  previous splits

Do not prune this tree

# Random forest


Example of decision frontier:





# Comparison (just an example)

[Book: The elements of statistical learning, Hastie, Tibshirani, Friedman, 2001]



## Pros

- Fast, parallelizable and appropriate for a large number of features
- Relatively easy to tune
- Frequently the winner in challenges

## Cons

- Overfitting if the size of the trees is too large
- Interpretability is lost (however importance of feature can be measured)

## Definition

A variable  $X^j$  is important to predict  $Y$  if breaking the link between  $X^j$  and  $Y$  increase the prediction error

$\{\bar{\mathcal{S}}_n^t = \mathcal{S}_n - \mathcal{S}_n^t, t = 1, \dots, n_{tree}\}$  **out-of-bag samples**: contains the samples not selected by bootstrap

# Variable importance

Let  $\{\bar{\mathcal{S}}_n^t = \mathcal{S}_n - \mathcal{S}_n^t, t = 1, \dots, n_{tree}\}$  **out-of-bag samples**


Let  $\{\bar{\mathcal{S}}_n^{t,j}, t = 1, \dots, n_{tree}\}$ : permuted out-of-bag-samples (the values of the  $j$ th variable have been randomly permuted).

$$\hat{I}(X^j) = \frac{1}{n_{tree}} \sum_{t=1}^{n_{tree}} R_n(h_t, \bar{\mathcal{S}}_n^{t,j}) - R_n(h_t, \bar{\mathcal{S}}_n^t)$$

with  $R_n(h, \mathcal{S})$ : empirical loss of  $h$  measured on  $\mathcal{S}$

# Variable importance: spam data

Spam dataset :



Motivation

Bagging

Random forests

Boosting

- AdaBoost as a Greedy Scheme

- Gradient Boosting

References

Motivation

Bagging

Random forests

Boosting

AdaBoost as a Greedy Scheme

Gradient Boosting

References

# A preliminary question

- **Is it possible to "boost" a weak learner into a strong learner ?**  
Michael Kearns
- Yoav Freund and Rob Schapire proposed an iterative scheme, called, Adaboost to solve this problem
  - **Idea**: train a sequence of learners on weighted datasets with weights depending on the loss obtained so far.
  - Freund and Schapire received the Godel prize in 2003 for their work on AdaBoost.




# Boosting a linear classifier

$$H_1(x) = h_1(x)$$

Binary Classifier:  $F_1(x) = \text{sign}(H_1(x))$

Here:  $h_1$ : linear classifier

Training error =  $R_n$




Source Jiri Matas (Oxford U.)

# Boosting a linear classifier


$$H_2(x) = \alpha_1 h_1(x) + \alpha_2 h_2(x)$$

Binary Classifier:  $F_2(x) = \text{sign}(H_2(x))$

$t = 2$




$R_n(H_t)$




Source Jiri Matas (Oxford U.)

# Boosting a linear classifier

$t = 3$




$R_n(H_t)$



Source Jiri

Matas (Oxford U.)

# Boosting a linear classifier




Source Jiri Matas (Oxford U.)

# Boosting a linear classifier

$t = 5$




$R_n(H_t)$




Source Jiri Matas (Oxford U.)

# Boosting a linear classifier

$t = 6$





$R_n(H_t)$



Source Jiri Matas (Oxford U.)

# Boosting a linear classifier

$t = 40$



Source Jiri Matas (Oxford U.)

**Definition: weak classifier**

A classifier whose average training error is no more than 0.5

NB : it means that we do not need to have a deep architecture as the base classifier (a "short" tree will fit for instance, a linear classifier will be perfect and so on...)



# Adaboost idea

1.  $\mathcal{H}$ : a chosen class of "weak" binary classifiers,  $\mathcal{A}$ : a learning algorithm for  $\mathcal{H}$ 
  - Set  $w_1(i) = 1/n$ ;  $H_0 = 0$
  - For  $t = 1$  to  $T$ 
    - $h_t = \arg \min_{h \in \mathcal{H}} \epsilon_t(h)$
    - with  $\epsilon_t(h) = \mathbb{P}_{i \sim \mathbf{w}_t}[h(x_i) \neq y_i]$
    - Choose  $\alpha_t$
    - Choose  $w_{t+1}$
    - $H_t = H_{t-1} + \alpha_t h_t$
  - Output  $F_T = \text{sign}(H_t)$

$\mathcal{H}$ : a chosen class of "weak" binary classifiers

- Set  $w_1(i) = 1/n$ ;  $H_0 = 0$
- For  $t = 1$  to  $T$ 
  - $h_t = \arg \min_{h \in \mathcal{H}} \sum_{i=1}^n \epsilon_t(h)$
  - With  $\epsilon_t(h) = \mathbb{P}_{i \sim \mathbf{w}_t}[h(x_i) \neq y_i]$
  - $\epsilon_t = \epsilon_t(h_t)$
  - $\alpha_t = \frac{1}{2} \log \frac{1-\epsilon_t}{\epsilon_t}$
  - let  $w_{t+1,i} = \frac{w_t(i)e^{-\alpha_t y_i h_t(x_i)}}{Z_{t+1}}$  where  $Z_{t+1}$  is a renormalization constant such that  $\sum_{i=1}^n w_{t+1,i} = 1$
- $H_t = H_{t-1} + \alpha_t h_t$

Output  $F_T = \text{sign}(H_T)$

# What weight to choose ?

With the chosen definition, we have:

$$\begin{aligned}w_{t+1,i} &= \frac{w_{t,i} e^{-\alpha_t y_i h_t(x_i)}}{Z_t} \\&= \frac{w_{t-1,i} e^{-\alpha_{t-1} y_i h_{t-1}(x_i)} e^{-\alpha_t y_i h_t(x_i)}}{Z_{t-1} Z_t} \\&= \frac{e^{-y_i \sum_{s=1}^t \alpha_s h_s(x_i)}}{n \prod_{s=1}^t Z_s} \\&= \frac{e^{-y_i H_t(x_i)}}{n \prod_{s=1}^t Z_s}\end{aligned}$$

You see the weights encourage to correct examples badly classified by the whole combination  $H_t$

## First of all let us study $Z_t$

$$\begin{aligned}Z_t &= \sum_{i=1}^n w_t(i) e^{-\alpha_t y_i h_t(x_i)} \\&= \sum_{i=1}^n w_t(i) e^{-\alpha_t y_i h_t(x_i)} \\&= \sum_{i: y_i h_t(x_i)=+1} w_t(i) e^{-\alpha_t} + \sum_{i: y_i h_t(x_i)=-1} w_t(i) e^{\alpha_t} \\&= (1 - \epsilon_t) e^{-\alpha_t} + \epsilon_t e^{\alpha_t} \\&= (1 - \epsilon_t) \sqrt{\frac{\epsilon_t}{1 - \epsilon_t}} + \epsilon_t \sqrt{\frac{1 - \epsilon_t}{\epsilon_t}} \\&= \dots \\&= 2\sqrt{\epsilon_t(1 - \epsilon_t)}\end{aligned}$$

## The training error theorem for boosting

The training error of the classifier returned by Adaboost at time  $T$  verifies:

$$R_n(F_T) \leq e^{-2 \sum_{t=1}^T (\frac{1}{2} - \epsilon_t)^2}.$$

Furthermore, if for all  $t \in [1, T]$ ,  $\gamma \leq (\frac{1}{2} - \epsilon_t)$ , then

$$R_n(F_T) \leq e^{-2\gamma^2 T}.$$

## Adaboost: Bound on the training error: proof

For all  $u \in \mathbb{R}$ , we have  $1_{u \leq 0} \leq \exp(-u)$ .

Then

$$\begin{aligned} R_n(F_T) &= \frac{1}{n} \sum_{i=1}^n 1_{y_i F_T(x_i) \leq 0} \\ &\leq \frac{1}{n} \sum_{i=1}^n \exp(-y_i F_T(x_i)) = \frac{1}{n} \sum_{i=1}^n \left[ n \prod_{t=1}^T Z_t \right]^{w_{t+1,i}} = \prod_{t=1}^T Z_t \end{aligned}$$

## Bound on the training error: proof ctd'

We can now express  $\prod Z_t$  in terms of  $\epsilon_t$ :

$$\begin{aligned}\prod_{t=1}^T Z_t &= \prod_{t=1}^T 2\sqrt{\epsilon_t(1-\epsilon_t)} \\ &= \text{by remarkable identity} \\ &= \prod_{t=1}^T \sqrt{1 - 4(1/2 - \epsilon_t)^2} \\ &\leq \prod_t e^{-2(1/2 - \epsilon_t)^2} = e^{-2 \sum_{t=1}^T (1/2 - \epsilon_t)^2}\end{aligned}$$

using the identity  $1 - u \leq \exp(-u)$ .

The proof reveals several interesting properties:

1.  $\alpha_t$  is chosen to minimize  $\prod_t Z_t = g(\alpha)$  with  $g(\alpha) = (1 - \epsilon_t)e^{-\alpha} + \epsilon_t e^{\alpha}$ 
  - $g'(\alpha) = -(1 - \epsilon_t)e^{-\alpha} + \epsilon_t e^{\alpha}$
  - $g'(\alpha) = 0$  iff  $(1 - \epsilon_t)e^{-\alpha} = \epsilon_t e^{\alpha}$  iff  $\alpha = 1/2 \log \frac{1-\epsilon_t}{\epsilon_t}$
2. The equality  $(1 - \epsilon_t)e^{-\alpha} = \epsilon_t e^{\alpha}$  means that Adaboost assigns at each time  $t$  the same distribution mass to correctly classified examples and incorrectly classified ones. However there is no contradiction because the number of incorrectly examples decreases.



<http://scikit-learn.org/stable/modules/ensemble.html#adaboost>

```
>>> from sklearn.crossvalidation import crossval_score
```

```
>>> from sklearn.datasets import load_iris
```

```
>>> from sklearn.ensemble import AdaBoostClassifier
```


```
>>> iris = load_iris()
```

```
>>> clf = AdaBoostClassifier(n_estimators=100)
```

```
>>> scores = crossval_score(clf, iris.data, iris.target)
```

```
>>> scores.mean() 0.9...
```

# Typical behavior of boosting



- You have to wait a long time to see Boosting overfit. However contrary to first assertions, Adaboost does overfit
- Early stopping: an answer
- or... bound with  $\ell_1$  norm the magnitude of the weights

Motivation

Bagging

Random forests

Boosting

AdaBoost as a Greedy Scheme

Gradient Boosting

References

At the same time, different groups proved that Adaboost writes as a coordinate descent in the convex hull of  $\mathcal{H}$ .

- Greedy function approximation, Friedman, 1999.
- MarginBoost and AnyBoost : Mason et al. 1999.

# Gradient Boosting: the idea

- At each boosting step, one need to solve

$$(h_t, \alpha_t) = \arg \min_{h, \alpha} \sum_{i=1}^n \ell(y_i, H_{t-1}(x_i) + \alpha h) = L(y, H_{t-1} + \alpha h)$$

- Gradient approximation  $L(y, H_{t-1} + \alpha h) \sim L(y, H_{t-1}) + \alpha \langle \nabla L(H_{t-1}), h \rangle$ .
- Gradient boosting: replace the minimization step by a *gradient descent* type step:
  - Choose  $h_t$  as the best possible descent direction in  $\mathcal{H}$
  - Choose  $\alpha_t$  that minimizes  $L(y, H + \alpha h_t)$
- Easy if finding the best descent direction is easy!

# Gradient boosting and Adaboost

Those two algorithms are equivalent!

- Denoting  $H_t = \sum_{t'=1}^t \alpha_{t'} h_{t'}$ ,

$$\begin{aligned} \sum_{i=1}^n e^{-y_i(H_{t-1}(x_i) + \alpha h(x_i))} &= \sum_{i=1}^n e^{-y_i H_{t-1}(x_i)} e^{-\alpha y_i h(x_i)} \\ &= \sum_{i=1}^n w'_i(t) e^{-\alpha y_i h(x_i)} \\ &= (e^\alpha - e^{-\alpha}) \sum_{i=1}^n w'_i(t) \ell^{0/1}(y_i, h(x_i)) \\ &\quad + e^{-\alpha} \sum_{i=1}^n w'_i(t) \end{aligned}$$

# Gradient boosting and adaboost (ctd)

Those two algorithms are equivalent!

- The minimizer  $h_t$  in  $h$  is independent of  $\alpha$  and is also the minimizer of

$$\sum_{i=1}^n w'_i(t) \ell^{0/1}(y_i, h(x_i))$$



# Gradient boosting and Adaboost

- The optimal  $\alpha_t$  is then given by

$$\alpha_t = \frac{1}{2} \log \frac{1 - \epsilon'_t}{\epsilon'_t}$$

with  $\epsilon'_t = (\sum_{i=1}^n w'_i(t) \ell^{0/1}(y_i, h_t(x_i))) / (\sum_{i=1}^n w'_i(t))$

- One verify then by recursion that

$$w_i(t) = w'_i(t) / \left( \sum_{i=1}^n w'_i(t) \right)$$

and thus the two procedures are equivalent!

# AnyBoost or Forward Stagewise Additive model

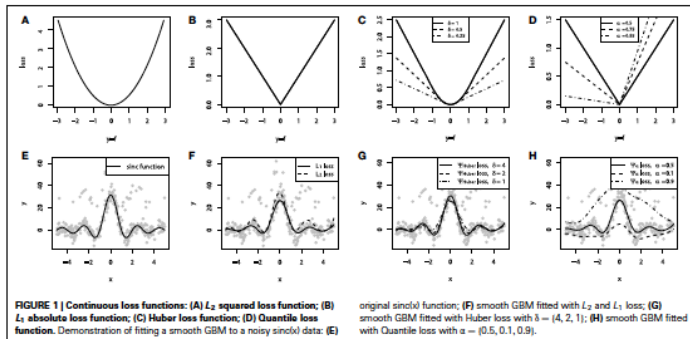
- General greedy optimization strategy to obtain a linear combination of *weak* predictor
  - Set  $t = 0$  and  $H_0 = 0$ .
  - For  $t = 1$  to  $T$ ,
    - $(h_t, \alpha_t) = \arg \min_{h, \alpha} \sum_{i=1}^n \ell(y_i, H_{t-1}(x_i) + \alpha h(x_i))$
    - $H_t = H_{t-1} + \alpha_t h_t$
  - Output  $H_T = \sum_{t=1}^T \alpha_t h_t$

# Losses in Forward Stagewise Additive Modeling

- AdaBoost with  $\ell(y, h) = e^{-yh}$
- LogitBoost with  $\ell(y, h) = \log(1 + e^{-yh})$
- $L_2$ Boost with  $\ell(y, h) = (y - h)^2$  (Matching pursuit)
- $L_1$ Boost with  $\ell(y, h) = |y - h|$
- HuberBoost with  $\ell(y, h) = |y - h|^2 \mathbf{1}_{|y-h| < \epsilon} + (2\epsilon|y - h| - \epsilon^2) \mathbf{1}_{|y-h| \geq \epsilon}$

Simple principle but no easy numerical scheme except for AdaBoost and  $L_2$ Boost...

# Continuous loss functions and gradient boosting



- Loss function for regression:  $\ell(y, h) = (y - h)^2$
- $(h_t, \alpha_t) = \arg \min_{h, \alpha} \sum_{i=1}^n (y_i - H_t(x_i) + \alpha h)^2$

Fitting the residuals.

Motivation

Bagging

Random forests

Boosting

References

# References

- Perrone, Cooper, When classifiers disagree, 1992
- Tumer and Gosh, 1996
- Breiman, Bagging predictors, 1996
- [Further reading](#): Buhlman and Yu, Analyzing bagging, Annals of stats., 2002
- Breiman, Random Forests, Machine Learning, 2001.
- Geurts, Ernst, Wehenkel, Extra-trees, JMLR, 2006
- Boosting, Freund and Shapire, 1996.
- Friedman, Gradient Boosting, 1999.